

**Institute_{of}
Data**

2022



Software Engineering

Module 10 - Part 2

Deployment And Maintenance



Institute of
Data

Agenda

Section 1 : Introduction to AWS EC2

Section 2 : Environment Variables

Section 3 : Deploying application in EC2

Section 4 : Beanstalk



Section 1 : Introduction to AWS EC2

Amazon EC2 offers a range of instance types that are tailored to certain use cases. Instance types are different combinations of CPU, memory, storage, and networking capabilities that allow you to choose the best resource mix for your applications. Each instance type has one or more instance sizes, allowing you to scale your resources to your target workload needs.



Compute using Amazon EC2

Whether you're building corporate, cloud-native, or mobile apps, or running enormous clusters to fuel analysis workloads, establishing and running your business starts with compute. AWS provides a comprehensive set of [compute services](#) that enable you to build, launch, run, and expand your applications and workloads on the world's most powerful, secure, and innovative cloud.

AWS computing services have the following characteristics:

- Right compute for your workloads
- Accelerate from idea to market
- Offer built-in security
- Flexibility to optimise costs
- Provide compute resource where you need it

AWS Compute



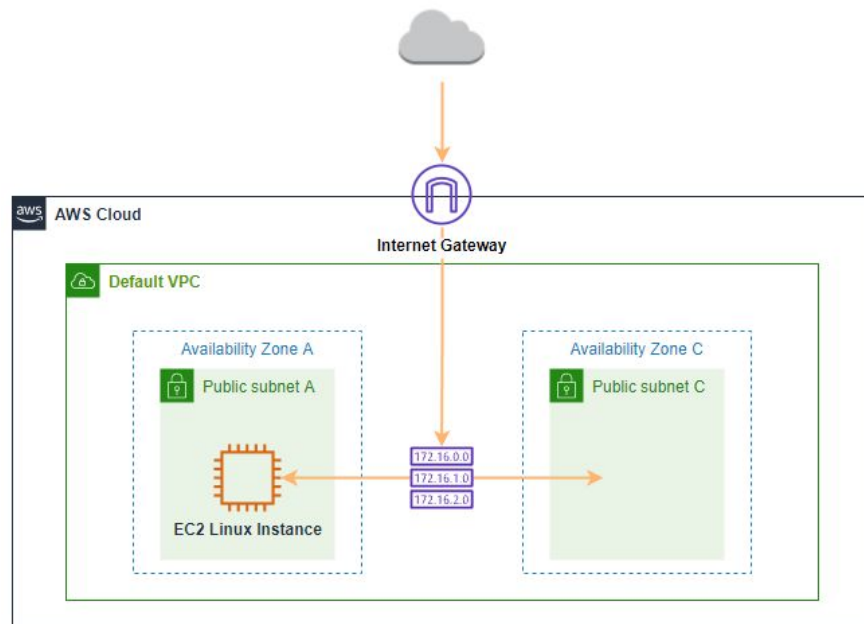


Overview

In the Amazon Web Services (AWS) Cloud, [Amazon EC2](#) delivers scalable computing capability. Using Amazon EC2 reduces the requirement for upfront hardware investment, allowing you to develop and deploy apps more quickly. Amazon EC2 allows you to create as many or as few virtual servers as you need, as well as establish security and networking and manage storage. You can scale up or down on Amazon EC2 to manage variations in demand or popularity spikes, decreasing the need to forecast traffic.

Create your own web server by going through the labs in the order below:

1. [Create a new key pair](#)
2. [Launch a Web Server Instance](#)
3. [Connect to your linux instance](#)
4. [Connect to EC2 Instance using PuTTY \(Optional\)](#)





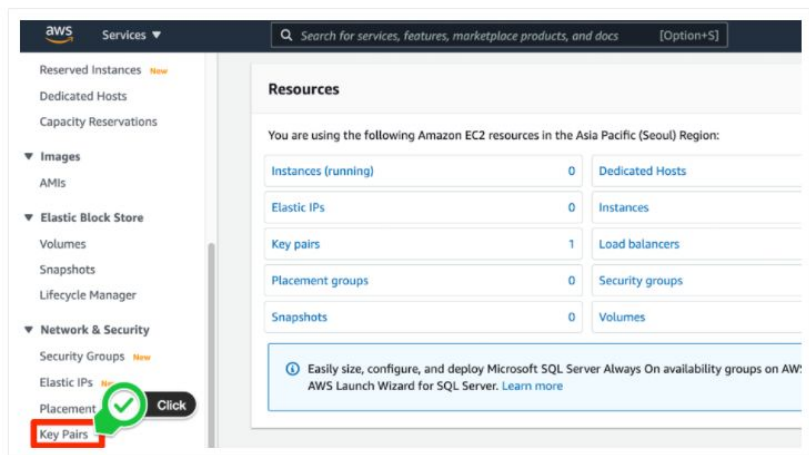
Create a key pair

In this lab, you will need to create an EC2 instance using an SSH keypair. The following steps outline creating a unique SSH keypair for you to use in this lab.

1. Sign into the AWS Management Console and open the [Amazon EC2 console](#). In the upper-right corner of the AWS Management Console, confirm you are in the desired AWS region.
2. Click on **Key Pairs** in the Network & Security section near the bottom of the leftmost menu. This will display a page to manage your SSH key pairs.
3. To create a new SSH key pair, click the **Create key pair** button at the top of the browser window.
4. Type **[Your Name]-ImmersionDay** into the Key Pair Name: text box and click **Create key pair** button.

*For Windows and PuTTY users, please select **ppk** for file format. The **pem** format should be used for Mac or SSH users.*

5. The page will download the file **[Your Name]-ImmersionDay.pem** (or ppk) to the local drive. Follow the browser instructions to save the file to the default download location. Remember the full path to the key pair file you just downloaded.



Create key pair

Key pair

A key pair, consisting of a private key and a public key, is a set of security credentials that you use to prove your identity when connecting to an instance.

Name

Type key pair name

AWS-ImmersionDay

The name can include up to 255 ASCII characters. It can't include leading or trailing spaces.

Key pair type

☒ RSA

☐ ED25519

Private key file format

☒ .pem

For use with OpenSSH

☐ .ppk

For use with PuTTY

Tags (Optional)

No tags associated with the resource.

Add tag

You can add 50 more tags.



Launch a Web server

We will launch an Amazon Linux 2 instance

1. Click on **EC2 Dashboard** near the top of the leftmost menu. And Click on **Launch instances**.

2. Enter a name for your instance and continue with the default Amazon Linux AMI (Amazon Machine Image) from the Quick Start tab:

The screenshot shows the AWS Management Console interface. On the left, the 'EC2 Dashboard' link is highlighted with a red box and a green checkmark, with a '1. Click' label. The main content area shows the 'Resources' section for the Asia Pacific (Seoul) Region, listing various EC2 resources. Below this, the 'Launch instance' section is visible, with a red box around the 'Launch instance' button and a green checkmark, labeled '2. Click'.

The screenshot shows the 'Name and tags' section of the AWS Management Console. The 'Name' field is highlighted with a red box and contains the text 'e.g. My Web Server'. Below this, the 'Application and OS Images (Amazon Machine Image)' section is visible, showing a search bar and a list of AMIs. The 'Amazon Linux 2 AMI (HVM) - Kernel 5.10, SSD Volume Type' is highlighted with a red box. The 'Quick Start' tab is selected, and the 'Amazon Linux' AMI is chosen from the 'Recents' list.



3. Next choose an *Instance Type*, select the **t2.micro** instance size which is eligible for the Free Tier

4. Next choose the key pair that you created back on slide 7. You will also need the local copy of this key to connect to your instance via PuTTY / SSH

The screenshot displays two sections of the AWS console. The first section, 'Instance type', shows the 't2.micro' instance type selected, which is 'Free tier eligible'. The second section, 'Key pair (login)', shows a dropdown menu with 'AWS-ImmersionDay' selected. A red box highlights the 't2.micro' instance type and the 'AWS-ImmersionDay' key pair. A blue box highlights the 'Key pair (login)' section header and the dropdown menu.

▼ Instance type [Info](#)

Instance type

t2.micro Free tier eligible

Family: t2 1 vCPU 1 GiB Memory

On-Demand Linux pricing: 0.0146 USD per Hour

On-Demand Windows pricing: 0.0192 USD per Hour

[Compare instance types](#)

▼ Key pair (login) [Info](#)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - required

Select

Q |

Proceed without a key pair (Not recommended) Default value

AWS-ImmersionDay

Type: rsa

AWS-ImmersionDay

JB-ImmersionDay

Type: rsa

vpc-03ad75a11be4d032b

[Create new key pair](#)

[Edit](#)



Network & Storage

- To ensure the right access is granted to your instance, make the following changes in Network settings:
 - **Allow SSH traffic** from My IP (*if your IP changes and you need to connect via SSH/PuTTY later on, remember to update this*)
 - **Tick Allow HTTPs traffic** from the internet
 - **Tick Allow HTTP traffic** from the internet
- Keep the default storage option as below

▼ **Configure storage** [Info](#) Advanced

1x 8 GiB gp2 ▼ Root volume

Free tier eligible customers can get up to 30 GB of EBS General Purpose (SSD) or Magnetic storage

Add new volume

0 x File systems Edit

▼ **Network settings** [Info](#) Edit

Network [Info](#)
vpc-03ad75a11be4d032b

Subnet [Info](#)
No preference (Default subnet in any availability zone)

Auto-assign public IP [Info](#)
Enable

Firewall (security groups) [Info](#)
A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

☒ Create security group ☐ Select existing security group

We'll create a new security group called 'launch-wizard-1' with the following rules:

☒ Allow SSH traffic from
Helps you connect to your instance

☐ Allow HTTPs traffic from the internet
To set up an endpoint, for example when c

☐ Allow HTTP traffic from the internet
To set up an endpoint, for example when c

My IP
49.177.92.48/32 ▲

Anywhere
0.0.0.0/0

Custom

My IP
49.177.92.48/32



Advanced Startup Script

7. In the Advanced tab, leave most settings unchanged, but scroll down to the **User data** section near the bottom. In here, paste the following startup script (taken from [here](#))

```
#!/bin/sh

# Install a LAMP stack
amazon-linux-extras install -y lamp-mariadb10.2-php7.2 php7.2
yum -y install httpd php-mbstring

# Start the web server
chkconfig httpd on
systemctl start httpd

# Install the web pages for our lab
if [ ! -f /var/www/html/immersion-day-app-php7.tar.gz ]; then
    cd /var/www/html
    wget
    https://aws-joozero.s3.ap-northeast-2.amazonaws.com/immersion-day-app-php7.tar.gz
    tar xvfz immersion-day-app-php7.tar.gz
fi

# Install the AWS SDK for PHP
if [ ! -f /var/www/html/aws.zip ]; then
    cd /var/www/html
    mkdir vendor
    cd vendor
    wget https://docs.aws.amazon.com/aws-sdk-php/v3/download/aws.zip
    unzip aws.zip
fi

# Update existing packages
yum -y update
```

▼ Advanced details Info

Purchasing option Info

☐ Request Spot Instances

Request Spot Instances at the Spot price, capped at the On-Demand price

Domain join directory Info

Select

Create new directory

User data Info

☐ User data has already been base64 encoded

8. Finally click 'Launch Instance' from the Summary on the right:

Cancel

Launch instance



✓ Success
Successfully initiated launch of instance (i-0ab2fd91be24faba3)

► Launch log

Next Steps

Get notified of estimated charges

Create [billing alerts](#) to get an email notification when estimated charges on your AWS bill exceed an amount you define (for example, if you exceed the free usage tier)

How to connect to your instance

Your instance is launching and it might be a few minutes until it is in the running state, when it will be ready for you to use

Click View Instances to monitor your instance's status. Once your instance is in the 'running' state, you can connect to it from the Instances screen. Find out [how to connect to your instance](#)

[View more resources to get you started](#)

[View all instances](#)

9. Your new instance will now launch, and will take a few minutes to initialise and start. Click **View all instances** to check progress.

10. Once your instance has launched, you will see your Web Server as well as the Availability Zone the instance is in, and the publicly routable **DNS name**. Click the checkbox next to your web server to view details about this EC2 instance.


The screenshot shows the AWS Management Console interface for the EC2 Instances page. The instance 'Web server for custom AMI' is highlighted with a red box. The instance is in the 'Running' state. The public IPv4 address is 52.78.23.186. The public IPv4 DNS is ec2-52-78-23-186.ap-northeast-2.compute.amazonaws.com. The private IPv4 address is 172.31.38.47. The private IPv4 DNS is ip-172-31-38-47.ap-northeast-2.compute.internal. The VPC ID is vpc-4a4ec021.



View your website live

1. Wait for the instance to pass the Status Checks to finish loading.
2. Open a new browser tab and browse the Web Server by entering the EC2 instance's **Public DNS name** into the browser. This name can be found in the console by reviewing the **Public IPv4 DNS** name line highlighted above. You should see a website that looks like the following.

If it doesn't load, try changing the address from https to http in the browser address bar.



LOAD TEST		RDS	
Meta-Data		Value	
Instanceld		i-0f9c0154bbc266ca9	
Availability Zone		ap-northeast-2c	

Current CPU Load: 1%



Connect to your Linux instance

1. In the EC2 instance console, select the instance you want to connect to, and then click the **Connect** button.
2. In the **Connect to instance** page, select **SSH client**. Follow the instructions provided.
3. If you are using Windows use PuTTY (next) or copy the ssh command from step 4.

Instances (1/1) info Connect Instance state Actions Launch instances

Filter instances

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability
Web server for custom AMI	i-0f9c0154bbc266ca9	Running	t2.micro	2/2 checks passed...	No alarms	ap-northeast-2

Connect to instance info

Connect to your instance i-0f9c0154bbc266ca9 (Web server for custom AMI) using any of these options

EC2 Instance Connect | **Session Manager** | **SSH client**

Instance ID
i-0f9c0154bbc266ca9 (Web server for custom AMI)

1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is AWS-ImmersionDay.pem
3. Run this command, if necessary, to ensure your key is not publicly viewable.
chmod 400 AWS-ImmersionDay.pem
4. Connect to your instance using its Public DNS:
ec2-52-78-23-186.ap-northeast-2.compute.amazonaws.com

Example:

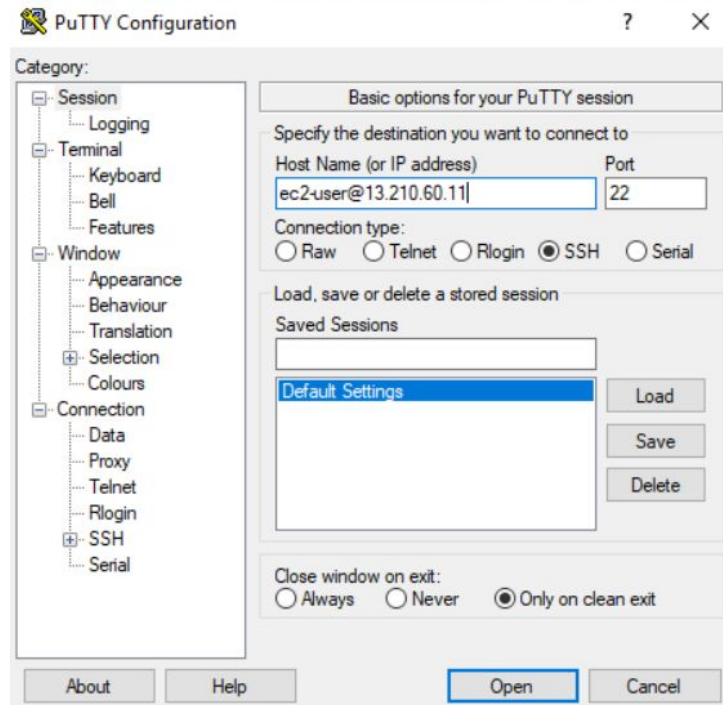
```
ssh -i "AWS-ImmersionDay.pem" ec2-user@ec2-52-78-23-186.ap-northeast-2.compute.amazonaws.com
```


Note: In most cases, the guessed user name is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI user name.



Connect using PuTTY (Windows)

1. Start PuTTY (if you need to download PuTTY first).
2. In the **Category** pane, choose **Session**.
3. In the **Host Name** box enter **ec2-user@[your public IP of EC2 that you created]**.
4. Set the **Port** value to 22.
5. Under **Connection type**, select **SSH**.
6. In the **Category** pane, expand Connection, expand **SSH**, and then choose **Auth**. Complete the following:
 - Choose **Browse**.
 - Select the **.ppk** file that you generated for your key pair
7. Go back to the Session tab, put a name in the box below Saved Sessions, and click Save. Next time you open PuTTY you can reuse this saved session.
8. Click **Open**. If this is the first time you have connected to this instance, PuTTY displays a security alert dialog box that asks whether you trust the host to which you are connecting. Choose **Yes**. A window opens and login as **ec2-user** and you are connected to your instance.





Exercise 1

Try creating your own EC2 ubuntu instance on AWS. Remember to create only a free tier server so that AWS does not charge you anything for your instance.



Section 2: Environment Variables

The majority of programmers believe environment variables to be key-value pairs supplied to a specific programme.

Both keys and values are always character sequences for each pair. You can also refer to them as strings. These pairs would be passed to a server application in backend projects.

However, front-end projects are not as simple. Browsers do not allow environment variables because they execute code. Additional libraries (plugins) are typically used by developers to replace the use of these variables with predefined constants. This bond usually occurs during the development of a project.



Environment Variables

The purpose of using environment variables is to keep setup and code separate. Modern products use the Infrastructure-as-Code model to manage infrastructure and specify variables as inextricable parts of that model. Why is it preferable to storing project configuration in JSON files?

For starters, the JSON format supports nested items. We try to keep the settings as simple as possible; nested information is not necessary. Second, we don't want to save credentials in project repositories directly. Finally, we want complete control over the settings we send to the application.



Dotenv NPM

Dotenv is a module that reads .env files and loads environment variables into process.env. It is one of the most used npm package for reading environment variables, and you have probably already used it in previous modules.

Firstly we need to install the package. For doing so we run the command

```
$ npm install dotenv --save
```

In order to use dotenv, simply add this code to the top of your JS file and you'll be able to use environment variables.

```
require('dotenv').config();
```



Dotenv NPM

To set these environment variables just create a .env file in the root directory of your project and add your secret tokens like this.

You can see here we added two environment variables and we utilised them in the server.js file and the dbConnect.js file of our application by using the dotenv npm, and on running the application we were able to use them in our respective files.

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL GITLENS: VISUAL FILE HISTORY

Navit@lessios-Mini mvc-structure % npm start

```
> mvc-structure@0.0.0 start
> node server.js
```

```
Listening on port 8000
MongoDB Connected: mongodb+srv://admin:VzC4x1HgCC2Fy1D6@cloudbootcamp.bv4zn.mongodb.net/myFirstDatabase?retryWrites=true&w=majority
```

```
.env
1 uri='mongodb+srv://admin:VzC4x1HgCC2Fy1D6@cloudbootcamp.bv4zn.mongodb.net/myFirstDatabase?retryWrites=true&w=majority'
2 PORT=8000

JS server.js > ...
You, 11 seconds ago | 1 author (You)
1 | require('dotenv').config() You, 11 seconds ago * Uncommitted changes
2 | let express = require("express");
3 | let app = express();
4 | let dbConnect = require("./dbConnect");
5
6 //dbConnect.dbConnect()
7 //var app = require('express')();
8 let http = require('http').createServer(app);
9 let io = require('socket.io')(http);
10 //const MongoClient = require('mongodb').MongoClient;
11
12 // routes
13 // let projectsRoute = require('./routes/projects')
14
15
16 var port = process.env.PORT || 8080;
17 app.use(express.json());
18 app.use(express.static(__dirname + '/public'));
19 let userRoute = require('./routes/userRoute')
20 app.use('/api/users',userRoute)
21
JS dbConnect.js > [e] mongooseOptions
6 | require('dotenv').config()
7 | var mongoose = require('mongoose');
8
9 const uri = process.env.uri || "mongodb://localhost/myFirstDatabase";
10
11
12 const mongooseOptions = {
13   useNewUrlParser: true, You, last month * Initial Commit ...
14   useUnifiedTopology: true
15 };
16 mongoose.set('useCreateIndex', true);
17 mongoose.set('useFindAndModify', false);
18
19 //Connect to MongoDB
20 mongoose.connect(uri, mongooseOptions, function (err) {
21   if (err) {
22     console.log("DB Error: ", err);
23     process.exit(1);
24   } else {
25     console.log('MongoDB Connected: ',process.env.uri);
26   }
27 });
28
```



Section 3: Deploying application in EC2

Once you have created your own ec2 instance from the previous slides you should be able to connect to it using ssh

```
~/ssh -- ubuntu@ip-172-31-41-169: ~ -- ssh -i aws_free_server_temp.pem ubuntu@54.252.135.104
```

```
Last login: Tue Apr 19 00:06:32 on ttys003
Navit@alessios-Mini ~ % cd .ssh
Navit@alessios-Mini .ssh % ssh -i aws_free_server_temp.pem ubuntu@54.252.135.104
Welcome to Ubuntu 20.04.3 LTS (GNU/Linux 5.11.0-1022-aws x86_64)
```

```
* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:        https://ubuntu.com/advantage
```

```
System information as of Mon Apr 18 14:08:08 UTC 2022
```

```
System load:  0.0           Processes:      103
Usage of /:   18.4% of 7.69GB Users logged in: 1
Memory usage: 20%          IPv4 address for eth0: 172.31.41.169
Swap usage:   0%
```

```
1 update can be applied immediately.
To see these additional updates run: apt list --upgradable
```

```
The list of available updates is more than a week old.
To check for new updates run: sudo apt update
```

```
Last login: Mon Apr 18 14:06:54 2022 from 220.253.142.233
To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.
```

```
ubuntu@ip-172-31-41-169:~$ █
```



Deploying application in EC2

Next what we need to do is to install Docker on our EC2 instance. The commands for this can vary depending on the type of linux installation used. If the left set of commands below are not working, try the right:

```
$ sudo apt update  
$ sudo apt install docker
```

```
$ sudo yum update  
$ sudo yum install docker
```

Once we have Docker installed on our server, we can check if it was successfully installed and start it by using the commands

```
$ docker --version  
$ sudo systemctl start docker.service
```

And we can also see if Docker is running or not by using the command

```
$ sudo systemctl status docker
```



Deploying application in EC2

Now that we have Docker installed all we need to do is to use the Docker image of our application that we created previously and hosted on Docker hub. First we need to pull the docker image on the EC2 instance.

```
$ sudo docker pull <image name>
```

```
$ sudo docker pull navitchoudhary22/mvc-structure
```

```
ubuntu@ip-172-31-41-169:~$ sudo docker pull navitchoudhary22/mvc-structure
Using default tag: latest
latest: Pulling from navitchoudhary22/mvc-structure
df9b9388f04a: Pull complete
622e2b598d8a: Pull complete
f7c8a32a53f2: Pull complete
7da04ed7d1ef: Pull complete
a7c1bda96431: Pull complete
b7d42cf4838a: Pull complete
3c8f7b92220d: Pull complete
62c017e40125: Pull complete
Digest: sha256:eb717e463f48bb92d50169a6e8ab496ecaf0b0a709c698d740d3c01a4bc476b1
Status: Downloaded newer image for navitchoudhary22/mvc-structure:latest
docker.io/navitchoudhary22/mvc-structure:latest
ubuntu@ip-172-31-41-169:~$
```

Next stop the existing application using port 80 by running the following command:

```
$ sudo systemctl stop httpd
```



Deploying application in EC2

So now we successfully cloned the image on our EC2 server, all we now need to do is to run the docker image

```
$ sudo docker run -d -p 80:8080 <image name>
```

```
$ sudo docker run -d -p 80:8080 navitchoudhary22/mvc-structure
```

```
~/ssh - ubuntu@ip-172-31-41-169: ~ -- ssh -i aws_free_server_temp.pem ubuntu@54.252.135.104
ubuntu@ip-172-31-41-169:~$ sudo docker run -d -p 8080:8080 navitchoudhary22/mvc-structure
c3207469f6dab7b365597f84bd2cab3b9f2a54682891ece2747d5551b467dbf8
ubuntu@ip-172-31-41-169:~$ docker ps
Got permission denied while trying to connect to the Docker daemon socket at unix:///var/run/docker.sock: Get "http://%2Fvar%2Frun%2Fdocker.sock/v1.24/containers/json": dial unix /var/run/docker.sock: connect: permission denied
ubuntu@ip-172-31-41-169:~$ sudo docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
c3207469f6da	navitchoudhary22/mvc-structure	"docker-entrypoint.s..."	13 seconds ago	Up 11 seconds	0.0.0.0:8080->8080/tcp, :::8080->8080/tcp	distracted_nightingale

```
ubuntu@ip-172-31-41-169:~$
```

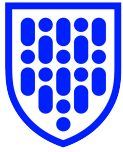
And now we have our application successfully running on our EC2 instance. You can see it running in your browser by using the public URL for your EC2 instance.

You can also see the demo application running [here](#).



Exercise 2

Try hosting the docker image of your nodejs application that you must have created using the CI/CD pipeline that you created previously. Share the link of your hosted application to your trainer.



Section 4: AWS Elastic Beanstalk

The entire application development process is being reshaped by cloud computing. A variety of cloud providers, such as Amazon Web Services and Microsoft Azure, provide development tools to make the process easier and more secure. The AWS Elastic Beanstalk development tool is an example of a PaaS-based development tool.

AWS Elastic Beanstalk is a simple tool for delivering and scaling web applications and services written in Java,.NET, PHP, Node.js, Python, Ruby, Go, and Docker on well-known servers like Apache, Nginx, Passenger, and IIS.

A developer can use AWS Elastic Beanstalk to launch an application without having to provision the underlying infrastructure while yet retaining high availability.





Benefits of Elastic Beanstalk

Offers Quicker Deployment: Elastic Beanstalk allows developers to quickly and simply deploy their apps. Users will not need to worry about the underlying infrastructure or resource settings because the application will be ready to use in minutes.

Supports Multi-Tenant Architecture: Customers can use AWS Elastic Beanstalk to distribute their programmes across numerous devices while ensuring scalability and security. It creates a detailed report on app usage and user profiles.

Simplifies Operations: Beanstalk is in charge of the application stack as well as the infrastructure provisioning and management. Developers must concentrate only on writing code for their application rather than managing and configuring servers, databases, firewalls, and networks.

Offers Complete Resource Control: Developers can use Beanstalk to select the appropriate AWS resources for their application, such as the EC2 instance type. It allows developers to have complete control over AWS resources and access them at any time.



Elastic Beanstalk Components

When deploying an application on Beanstalk there are certain terms that will come up frequently. Let us look at those concepts:

Application:

- In Elastic Beanstalk, an application is conceptually comparable to a folder.
- An application is made up of various components such as environments, versions, and configurations.

Application Version:

- A specific, identified iteration of deployable code for a web application is referred to as an application version.
- An Amazon S3 object containing deployable code, such as a Java WAR file, is referenced by an application version.

Environment:

- The current version of the Elastic Beanstalk Application will be active in environments within the Elastic Beanstalk Application.
- At any one time, each environment only runs one application version. However, the same or different versions of an application can be operated in many settings at the same time.

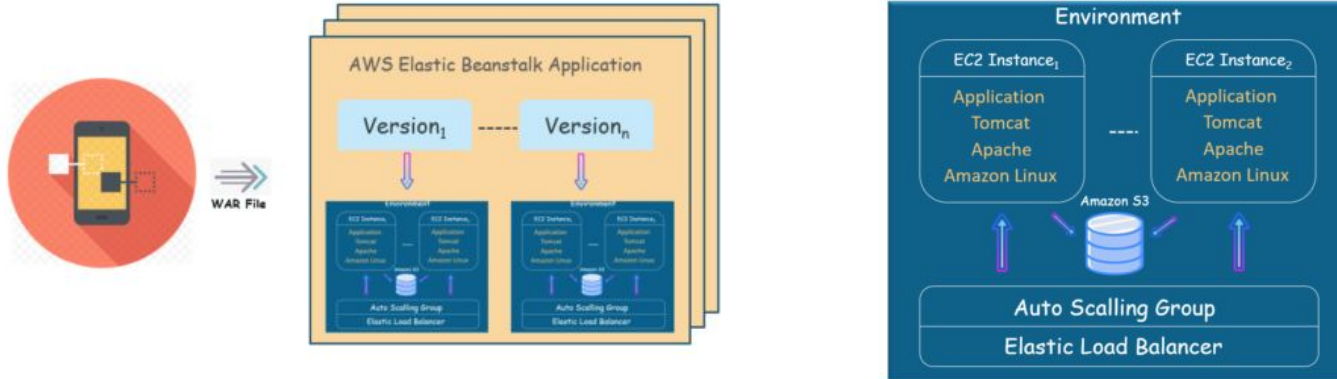


Elastic Beanstalk Components

Environment Tier:

Based on requirements Beanstalk offers two different Environment tiers: Web Server Environment, Worker Environment

- Web Server Environment: Handles HTTP requests from clients
- Worker Environment: Processes background tasks which are resource consuming and time intensive





Setting up Elastic Beanstalk Server

1. Our first step is to setup Elastic Beanstalk server on Amazon AWS. Login to your AWS Management Console and click on “Elastic Beanstalk” under services (or search for it)
2. Next step is to create an Amazon Elastic Beanstalk application. Click on the “Create Application” button and choose a **web server environment**.

In the create application page fill out the details as shown below (leave the rest as default):

Application Name: Hello-Express

Platform: Node.js

Platform branch: Node.js 16 running on 64bit Amazon Linux 2

Platform version: 5.6.2 (Recommended)

Application Code: Sample Code

Finally, click on the “Create application” button. This will take few minutes to process and setup the application.

AWS services

▼ All services

Compute

EC2

Lightsail

Lambda

Batch

Elastic Beanstalk

Serverless Application

Repository

AWS Outposts

EC2 Image Builder

AWS App Runner

Quantum Technologies

Amazon Braket

Management & Governance

AWS Organizations

CloudWatch

AWS Auto Scaling

CloudFormation

CloudTrail

Config

OpsWorks

Service Catalog

Security, Identity, & Compliance

IAM

Resource Access Manager

Cognito

Secrets Manager

GuardDuty

Inspector

Amazon Macie

AWS Single Sign-On

Certificate Manager

Key Management Service

Compute

Amazon Elastic Beanstalk

End-to-end web application management.

Amazon Elastic Beanstalk is an easy-to-use service for deploying and scaling web applications and services developed with Java, .NET, PHP, Node.js, Python, Ruby, Go, and Docker on familiar servers such as Apache, Nginx, Passenger, and IIS.

Get started

Easily deploy your web application in minutes.

[Create Application](#)

Elastic Beanstalk > Environments > Helloexpress-env



Creating Helloexpress-env

This will take a few minutes.

9:49pm Using elasticbeanstalk-us-east-1-200057541580 as Amazon S3 storage bucket for environment data.

9:49pm createEnvironment is starting.



Setting up Elastic Beanstalk Server

Once the create application setup is finished you should see a similar screen as shown in the screenshot.

The screenshot above indicates that the application “Hello-Express” has been created and it also contains a default environment “Helloexpress-env”.

Since we want our GitHub changes to propagate and deployed to the AWS Elastic Beanstalk, we must setup a Pipeline on AWS management console.

Search for ‘CodePipeline’ in the top AWS search bar.

Environment name	Health	Application name	Date created	Last modified	URL	Running versions	Platform	Platform state	Tier name
Helloexpress-env	OK	Hello-Express	2021-07-26 21:49:12 UTC-0500	2021-07-26 21:51:35 UTC-0500	Helloexpress-env-ea-pd2qx38s.us-east-1.elasticbeanstalk.com	Sample Application	Node.js 14 running on 64bit Amazon Linux 2	Supported	WebServer



Developer Tools

CodeStar

CodeCommit

CodeArtifact

CodeBuild

CodeDeploy

CodePipeline

Cloud9

CloudShell

X-Ray

AWS FIS

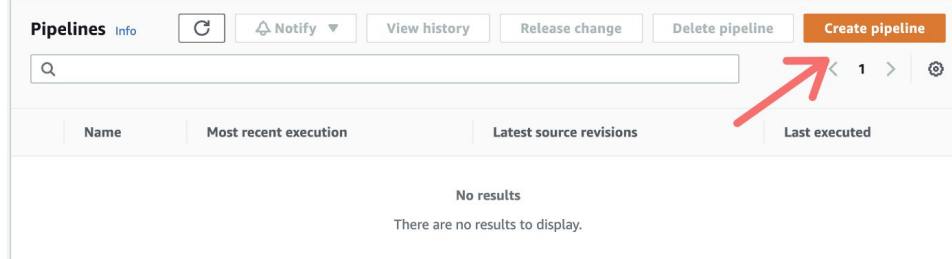


Setting up Elastic Beanstalk Server

This will open up a screen allowing you to create a new pipeline. Click the “Create pipeline” button.

After clicking on the “Create pipeline” button you will be taken to a page, where you can add details about the pipeline.

Once you fill out the “Pipeline name”, it will automatically fill out the Role name. Click “Next” to continue.



The screenshot shows the AWS CodePipeline console. At the top, there's a navigation bar with 'Pipelines' and 'Info' tabs. Below the navigation bar, there's a search bar and a table with columns: 'Name', 'Most recent execution', 'Latest source revisions', and 'Last executed'. The table is currently empty, displaying 'No results' and 'There are no results to display.' A red arrow points to the 'Create pipeline' button in the top right corner.

Pipeline settings

Pipeline name
Enter the pipeline name. You cannot edit the pipeline name after it is created.

No more than 100 characters

Service role

☒ **New service role**
Create a service role in your account

☐ **Existing service role**
Choose an existing service role from your account

Role name

Type your service role name

☒ **Allow AWS CodePipeline to create a service role so it can be used with this new pipeline**

Advanced settings

Cancel Next



Setting up Elastic Beanstalk Server

In the next screen, you will choose the source stage. Since, we are using GitHub as our source repository we will choose GitHub (Version 2).

GitHub (Version 1) is no longer recommended as it does not use the updated authentication methods

When you select GitHub (Version 2) an additional form will show up which will allow you to integrate your GitHub repository to the pipeline. Click on “Connect to GitHub” to start the process.

Add source stage [Info](#)

Source

Source provider

This is where you stored your input artifacts for your pipeline. Choose the provider and then provide the connection details.

GitHub (Version 2)



New GitHub version 2 (app-based) action

To add a GitHub version 2 action in CodePipeline, you create a connection, which uses GitHub Apps to access your repository. Use the options below to choose an existing connection or create a new one. [Learn more](#)

Connection

Choose an existing connection that you have already configured, or create a new one and then return to this task.



or

Connect to GitHub

Repository name

Choose a repository in your GitHub account.



No options

Branch name

Choose a branch of the repository.



Change detection options

☒ Start the pipeline on source code change

Automatically starts your pipeline when a change occurs in the source code. If turned off, your pipeline only runs if you start it manually or on a schedule.

Output artifact format

Choose the output artifact format.



CodePipeline default

AWS CodePipeline uses the default zip format for artifacts in the pipeline. Does not include git metadata about the repository.



Full clone

AWS CodePipeline passes metadata about the repository that allows subsequent actions to do a full git clone. Only supported for AWS CodeBuild actions.



Setting up Elastic Beanstalk Server

When you click “Connect to GitHub” it will open a small popup, which will allow you to create a connection. Add a connection name and click on the “Connect to GitHub” button.

You will have to enter your GitHub credentials to create a connection between GitHub and Code pipeline. Once the connection has been made, it will allow you to pick your GitHub repository and the branch

We have selected “Hello-Express” repository, which was created earlier (you can choose any standalone Node.js app) and the master branch of the repository because we want to deploy the code in the master branch to the server.

Click next.

Add source stage [Info](#)

Source

Source provider

This is where you stored your input artifacts for your pipeline. Choose the provider and then provide the connection details.

GitHub (Version 2)



New GitHub version 2 (app-based) action

To add a GitHub version 2 action in CodePipeline, you create a connection, which uses GitHub Apps to access your repository. Use the options below to choose an existing connection or create a new one. [Learn more](#)

Connection

Choose an existing connection that you have already configured, or create a new one and then return to this task.

arn:aws:codestar-connections:us-east-1:200057541580:connection/f43f2ce3

[Connect to GitHub](#)



Ready to connect

Your GitHub connection is ready for use.

Repository name

Choose a repository in your GitHub account.

azamsharp/Hello-Express

<account>/<repository-name>

Branch name

Choose a branch of the repository.

master

Change detection options

☒ Start the pipeline on source code change

Automatically starts your pipeline when a change occurs in the source code. If turned off, your pipeline only runs if you start it manually or on a schedule.

Output artifact format

Choose the output artifact format.



CodePipeline default

AWS CodePipeline uses the default zip format for artifacts in the pipeline. Does not include git metadata about the repository.



Full clone

AWS CodePipeline passes metadata about the repository that allows subsequent actions to do a full git clone. Only supported for AWS CodeBuild actions.

Cancel

Previous

Next



Setting up Elastic Beanstalk Server

The next screen will allow you to add a build stage. We are going to skip this step, so click on the “Skip build stage” button at the bottom.

A confirmation dialog will popup, select “Skip”

Add build stage Info

Build - optional

Build provider

This is the tool of your build project. Provide build artifact details like operating system, build spec file, and output file names.

AWS CodeBuild

Region

US East (N. Virginia)

Project name

Choose a build project that you have already created in the AWS CodeBuild console. Or create a build project in the AWS CodeBuild console and then return to this task.

Q

or

[Creating project](#)

Environment variables - optional

Choose the key, value, and type for your CodeBuild environment variables. In the value field, you can reference variables generated by CodePipeline. [Learn more](#)

Add environment variable

Build type



Single build

Triggers a single build.



Batch build

Triggers multiple builds as a single execution.

Cancel

Previous

Skip build stage

Next





Setting up Elastic Beanstalk Server

Next we will land on the “Add deploy stage” page. This is where we need to select our AWS Elastic Beanstalk application and its environment.

Click next.

The next screen will be the review screen. Make sure all settings and configurations are correct. Scroll down at the bottom and click “Create pipeline” button.

Add deploy stage [Info](#)



You cannot skip this stage

Pipelines must have at least two stages. Your second stage must be either a build or deployment stage. Choose a provider for either the build stage or deployment stage.

Deploy

Deploy provider

Choose how you deploy to instances. Choose the provider, and then provide the configuration details for that provider.

AWS Elastic Beanstalk

Region

US East (N. Virginia)

Application name

Choose an application that you have already created in the AWS Elastic Beanstalk console. Or create an application in the AWS Elastic Beanstalk console and then return to this task.

Hello-Express

Environment name

Choose an environment that you have already created in the AWS Elastic Beanstalk console. Or create an environment in the AWS Elastic Beanstalk console and then return to this task.

Helloexpress-env

Cancel

Previous

Next

Step 3: Add build stage

Build action provider

Build stage
No build

Step 4: Add deploy stage

Deploy action provider

Deploy action provider
AWS Elastic Beanstalk
ApplicationName
Hello-Express
EnvironmentName
Helloexpress-env

Cancel

Previous

Create pipeline



Setting up Elastic Beanstalk Server

As soon as the pipeline is setup, it will try to deploy the app to Amazon AWS Elastic Beanstalk.

After a while it comes back with deployment succeeded status.

Developer Tools > CodePipeline > Pipelines > Hello-Express-Pipeline

Hello-Express-Pipeline [Notify](#) [Edit](#) [Stop execution](#) [Clone pipeline](#) [Release change](#)

Source Succeeded
Pipeline execution ID: 143c07e3-084f-4b8c-abe5-2eb96de4a24d

Source
GitHub (Version 2) [Info](#)
Succeeded - 1 minute ago
4028e681 [View](#)

4028e681 [View](#) Source: adding hello express

[Disable transition](#)

Deploy In progress
Pipeline execution ID: 143c07e3-084f-4b8c-abe5-2eb96de4a24d

Deploy
AWS Elastic Beanstalk [Info](#)
In progress - 1 minute ago

4028e681 [View](#) Source: adding hello express

Developer Tools > CodePipeline > Pipelines > Hello-Express-Pipeline

Hello-Express-Pipeline [Notify](#) [Edit](#) [Stop execution](#) [Clone pipeline](#) [Release change](#)

Source Succeeded
Pipeline execution ID: 10126dcb-f260-4c71-affa-5c247d79b634

Source
GitHub (Version 2) [Info](#)
Succeeded - 2 minutes ago
96e443c2 [View](#)

96e443c2 [View](#) Source: Update app.js

[Disable transition](#)

Deploy Succeeded
Pipeline execution ID: 10126dcb-f260-4c71-affa-5c247d79b634

Deploy
AWS Elastic Beanstalk [Info](#)
Succeeded - Just now

96e443c2 [View](#) Source: Update app.js



Add Environment Variables

You may also need to set the internal port for your application. To do this, go to your Beanstalk Environment and click on Configuration in the left hand menu, then click on Edit next to Software:

Scroll down to the Environment Properties section and enter the default port used by your application (check server.js or index.js):

Click Apply Changes when done.

Elastic Beanstalk > Environments > Socketchat-env > Configuration

Configurations

Cancel Review changes Apply all changes

Table View

Search for an option name or value

Category	Options	Actions
Software	Log streaming: disabled Proxy server: nginx Rotate logs: disabled X-Ray daemon: disabled	Edit

Environment properties

The following properties are passed in the application as environment properties. [Learn more](#)

Name	Value
PORT	3000



Setting up Elastic Beanstalk Server

Now we need to see if we can access our routes or not. Click on “All services” and then select Elastic Beanstalk

Next click on the URL of your Hello-Express application. This will launch the root endpoint.

You can also click on the environment name and it will take you to a page that gives you all the details about the application and also lets you see the logs, monitor the application and many more.

Now you have successfully created an application on AWS Elastic Beanstalk.

The screenshot shows the 'All environments' page in the AWS Elastic Beanstalk console. At the top, there's a search bar and a 'Create a new environment' button. Below is a table with columns: Environment name, Health, Application name, Date created, Last modified, URL, and Running versions. The first row shows an environment named 'Helloexpress-env' with a health status of 'Ok', application 'Hello-Express', and a URL 'Helloexpress-env.eba-pd2qx38s.us-east-1.elasticbeanstalk.com'. A red arrow points to this URL.

Environment name	Health	Application name	Date created	Last modified	URL	Running versions
Helloexpress-env	Ok	Hello-Express	2021-07-26 21:49:12 UTC-0500	2021-07-27 21:25:37 UTC-0500	Helloexpress-env.eba-pd2qx38s.us-east-1.elasticbeanstalk.com	code-pipeline-1627496e443c21eef71126

The screenshot shows the details page for the 'Helloexpress-env' environment. The left sidebar has a 'Logs' link highlighted with a red box. The main content area shows the environment's health as 'Ok', its running version, and the platform (Node.js 14). There's a 'Recent events' section at the bottom showing a log deletion event.

Environments
Applications
Change history

▼ Hello-Express
Application versions
Saved configurations

▼ Helloexpress-env
Go to environment
Configuration
Logs
Health
Monitoring
Alarms
Managed updates
Events
Tags

Helloexpress-env
Helloexpress-env.eba-pd2qx38s.us-east-1.elasticbeanstalk.com
Application name: Hello-Express

Refresh Actions

Health
Ok
Causes

Running version
code-pipeline-1627496e443c21eef7112612633100991fac4b384a81f
Upload and deploy

Platform
nodejs
Node.js 14 running on 64bit Amazon Linux 2/5.4.3
Change

Recent events
Show all

Time	Type	Details
2021-07-27 21:28:47 UTC-0500	INFO	Deleted log fragments for this environment.



Exercise 3

Try hosting the GitHub repo of the nodejs application that you must have created using AWS Elastic Beanstalk. Share the link of your hosted application with your trainer.



Other Deployment options

We can also look into a few other deployment options for deploying our application on the cloud. Some other popular options would be

- [IBM Cloud Engine](#) (Paid Service)
- [Kubernetes](#) (Paid and High difficulty level)

End of Presentation